


City-GAN: a Conditional GAN Application

by Anas Hamed



About the Paper

- City-GAN: Learning architectural styles using a custom Conditional GAN architecture^[1]
- Maximilian Bachl & Daniel C. Ferreira
- Published 3 July 2019
- <https://github.com/muxamilian/city-gan>



City-GAN Overview

Aims to

- Capture architectural features of certain cities

by

- Training on pictures of facades
- Generating new ones in a specific style

Why this paper was chosen

- Builds upon DC-GANs^[2] and Conditional GANs^[3]
- Outlines the limitations and shortcomings of such methods
- Showcases recent applications and developments
- Introduces new ideas

The dataset


- Google Street View images
- Four cities: Amsterdam, Manhattan, Paris, Vienna
- 1000 Images per city were chosen such that
 - Facades are fully visible
 - No anomalous structures (no construction sites or occlusions)

Examples of anomalous samples (not used)




Examples of valid samples

Amsterdam




Manhattan



Paris



Vienna





Preprocessing & Augmentation

640x640 pixels



Stochastic flip




Random crop



- 64x64
- 128x128
- 256x256

First Approach: Standard GAN

- Deep Convolutional GAN^[2]
- No distinguishing factors between different cities
- Discriminator: continuously convolute input image
- Generator: continuously deconvolute noise vector



DC-GAN: Results

- Mode collapse
- Images too blurry
- Facades have no specific style
- Architecture does not adequately represent features

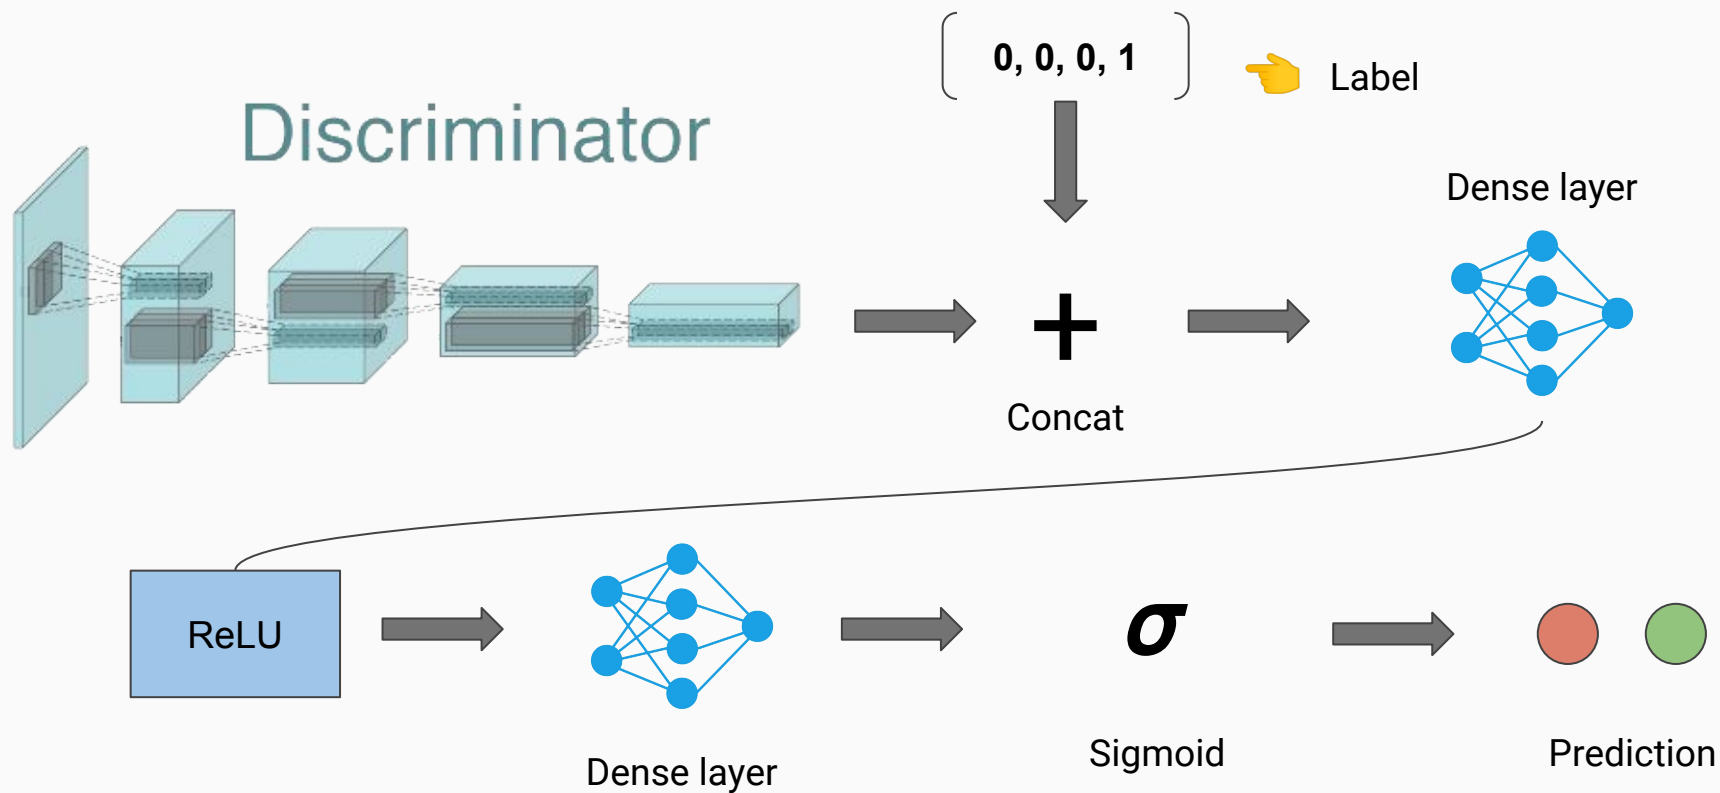
DC-GAN: Results



Second Approach: Conditional GAN

- Motive: allow Gan to focus on one style at a time
- Generator
 - Same architecture as DC-GAN
 - Concatenate label with latent variable
- Discriminator
 - Concatenate labels with output
 - Add dense layers

C-GAN Discriminator Architecture




C-GAN: Results

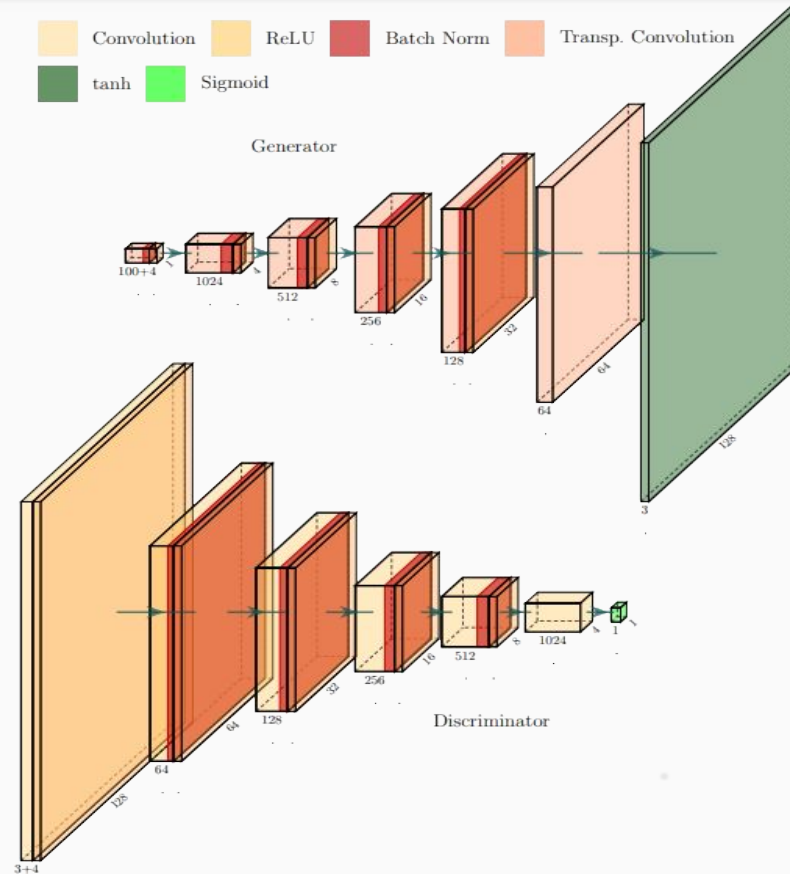


City-GAN approach


- Generator: concatenate label with latent variable
- Discriminator: concatenate label with each pixel



City-GAN Architecture



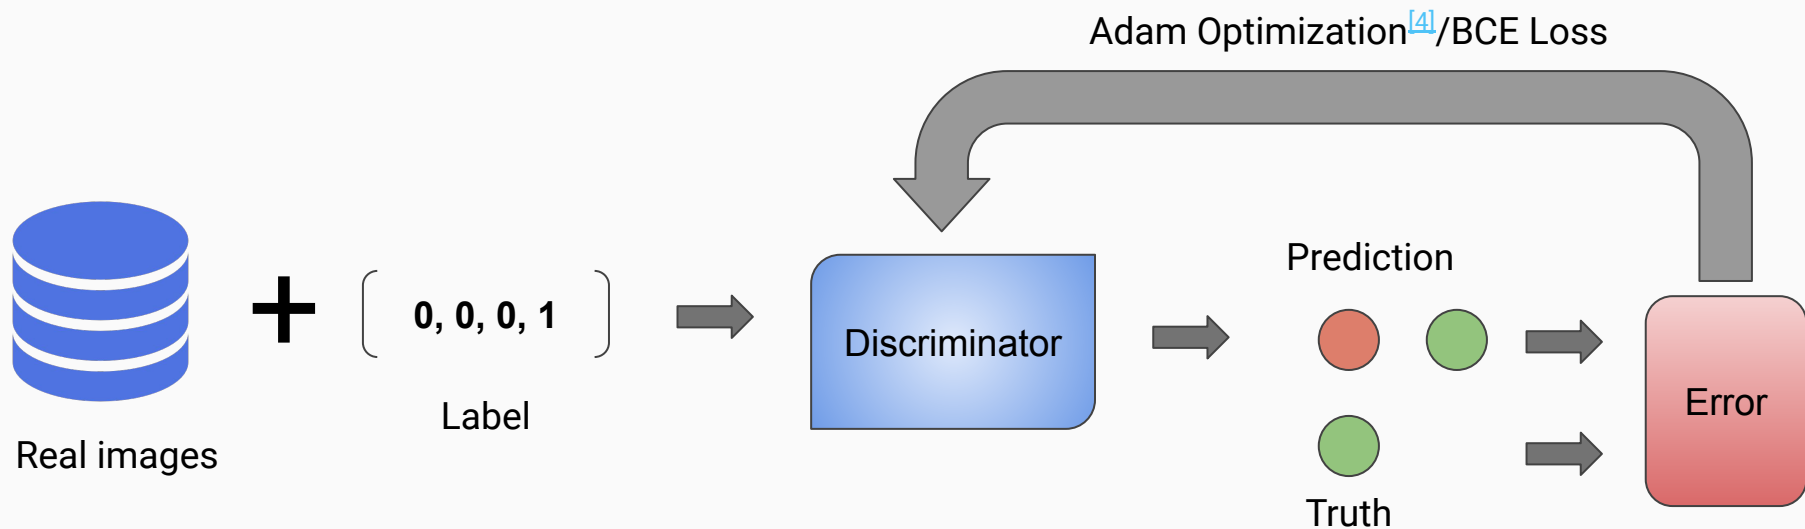
City-GAN: Results



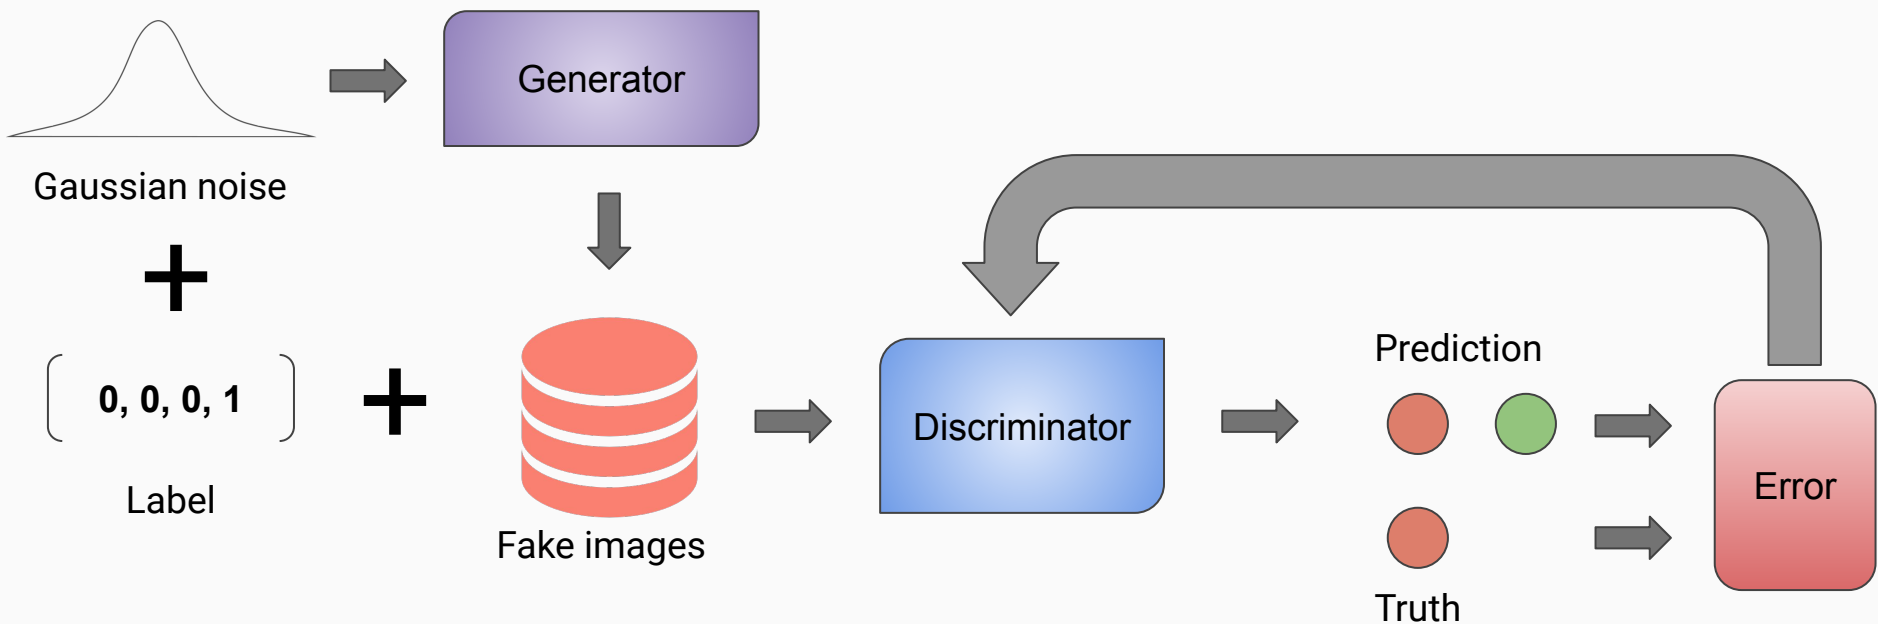
Training

- Discriminator
 - Real images (1 batch)
 - Fake images (1 batch)
- Generator
 - Produce images, adjust weights based on discriminator feedback

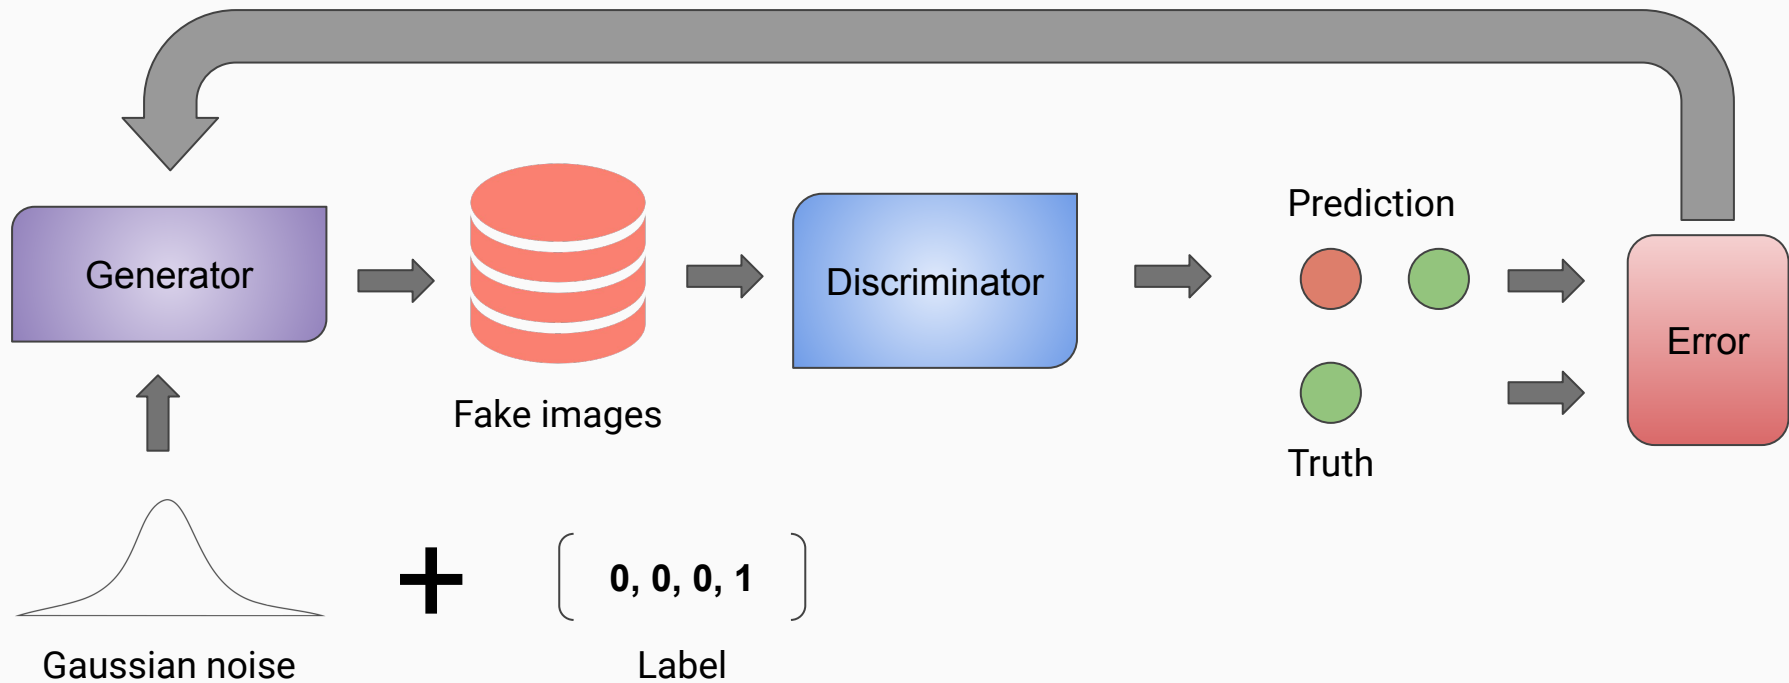
Discriminator Training: Phase 1



Discriminator Training: Phase 2



Generator Training



GAN Training

```
1 for batch in batches:
2     ##### DISCRIMINATOR TRAINING: PHASE 1 #####
3     batch_with_labels = concat_with_each_pixel(class_one_hot, batch)
4     train_discriminator(batch_with_labels, ground_truth=np.ones((batch_size, 1)))
5
6     ##### DISCRIMINATOR TRAINING: PHASE 2 #####
7     noise = torch.randn(batch_size, latent_dim) # Gaussian noise
8     noise_with_labels = np.concatenate((class_one_hot, noise), axis=1) # Concat class with each latent vector
9     fake_images = generator.generate(input=noise_with_labels, size=batch_size)
10    fake_images_with_labels = concat_with_each_pixel(class_one_hot, fake_images)
11    train_discriminator(fake_images_with_labels, ground_truth=np.zeros((batch_size, 1)))
12
13    ##### GENERATOR TRAINING #####
14    fake_preds = discriminator.evaluate(fake_images_with_labels)
15    loss_function = torch.nn.BCELoss() # Binary Cross Entropy loss
16    error = loss_function(output=fake_preds, ground_truth=np.ones((batch_size, 1)))
17    gradient_generator = error.backward() # Computes the gradient
18    # Adam optimizer
19    optimizer = torch.optim.Adam(generator, error, gradient_generator, lr=0.0002, betas=(0.5, 0.999))
20    optimizer.step()
21
22 def train_discriminator(images, ground_truth):
23     image_preds = discriminator.evaluate(images)
24     loss_function = torch.nn.BCELoss() # Binary Cross Entropy loss
25     error = loss_function(output=image_preds, ground_truth=ground_truth)
26     gradient_discriminator = error.backward() # Computes the gradient
27     # Adam optimizer
28     optimizer = torch.optim.Adam(discriminator, error, gradient_discriminator, lr=0.0002, betas=(0.5, 0.999))
29     optimizer.step()
```


Generator Layers

```
Generator(  
  (main): Sequential(  
    (0): ConvTranspose2d(104, 1024, kernel_size=(4, 4), stride=(1, 1), bias=False)  
    (1): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (2): ReLU(inplace=True)  
  
    (3): ConvTranspose2d(1024, 512, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)  
    (4): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (5): ReLU(inplace=True)  
  
    (6): ConvTranspose2d(512, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)  
    (7): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (8): ReLU(inplace=True)  
  
    (9): ConvTranspose2d(256, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)  
    (10): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (11): ReLU(inplace=True)  
  
    (12): ConvTranspose2d(128, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)  
    (13): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (14): ReLU(inplace=True)  
  
    (15): ConvTranspose2d(64, 3, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)  
    (16): Tanh()  
  )  
)
```


Discriminator Layers

```
Discriminator(  
  (main): Sequential(  
    (0): Conv2d(7, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)  
    (1): LeakyReLU(negative_slope=0.2, inplace=True)  
  
    (2): Conv2d(64, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)  
    (3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (4): LeakyReLU(negative_slope=0.2, inplace=True)  
  
    (5): Conv2d(128, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)  
    (6): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (7): LeakyReLU(negative_slope=0.2, inplace=True)  
  
    (8): Conv2d(256, 512, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)  
    (9): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (10): LeakyReLU(negative_slope=0.2, inplace=True)  
  
    (11): Conv2d(512, 1024, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)  
    (12): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (13): LeakyReLU(negative_slope=0.2, inplace=True)  
  
    (14): Conv2d(1024, 1, kernel_size=(4, 4), stride=(1, 1), bias=False)  
    (15): Sigmoid()  
  )  
(end): Sequential()  
)
```


Effect of Image Size

Varying levels of success were observed for different image sizes:

- **64x64:** best results
- **128x128:** some artifacts but generally acceptable
- **256x256:** no satisfactory results



Additional Dataset

- Larger, uncurated dataset
- Stanford university
- Washington DC, Las Vegas
- Same quality as curated dataset



Interpolation



Amsterdam → Florence



Amsterdam → D.C.



Amsterdam → Manhattan



Europe → U.S.

Label Weights

- Sign and magnitude control feature presence

Amsterdam	Florence	DC	Manhattan
1	0	0	0
1/2	1/2	0	0
0	0	1/2	1/2
1	0	0	-1

Amsterdam

Europe

U.S.

Amsterdam - Manhattan

Label Weights

- Amsterdam - Manhattan: $(1, 0, 0, 0) - (0, 0, 0, 1) = (1, 0, 0, -1)$



- Europe \rightarrow U.S.: $(1/2, 1/2, 0, 0) \rightarrow (0, 0, 1/2, 1/2)$



Wrap-up

- DC-GAN not suited to the task
 - Conditional GANs did not produce satisfactory results
 - Adding labels to images produced better results
 - Larger image sizes led to non-convergence
-
- Thoughts/questions?

References

- [1] Maximilian Bachl and Daniel Ferreira. 2019. City-GAN: Learning architectural styles using a custom Conditional GAN architecture. arXiv:1907.05280 [cs.CV] (Jul. 2019). <https://arxiv.org/abs/1907.05280> arXiv:1907.05280.
- [2] Alec Radford, Luke Metz, and Soumith Chintala. 2015. Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks. arXiv:1511.06434 [cs] (Nov. 2015). <http://arxiv.org/abs/1511.06434> arXiv: 1511.06434.
- [3] Mehdi Mirza and Simon Osindero. 2014. Conditional Generative Adversarial Nets. arXiv:1411.1784 [cs, stat] (Nov. 2014). <http://arxiv.org/abs/1411.1784> arXiv: 1411.1784.
- [4] Diederik Kingma and Jimmy Lei Ba. 2017. Adam: a method for Stochastic Optimization. arXiv:1412.6980 [cs.LG] (Jan. 2017). <https://arxiv.org/abs/1412.6980> arXiv:1412.6980v9.

Thanks!