

IMAGE-TO-IMAGE TRANSLATION

APPLICATIONS OF CYCLEGAN

Prepared for:

CSE 728 Presentation



Introduction

- CycleGAN
- Applications

Methods

- Unpaired Data Mapping
- Model Architecture
- Cycle Consistency Loss
- Generator and Discriminator
- Loss functions
- Optimizer
- Model Training

Discussion

- Broader Application
- Limitation
- Code Demo

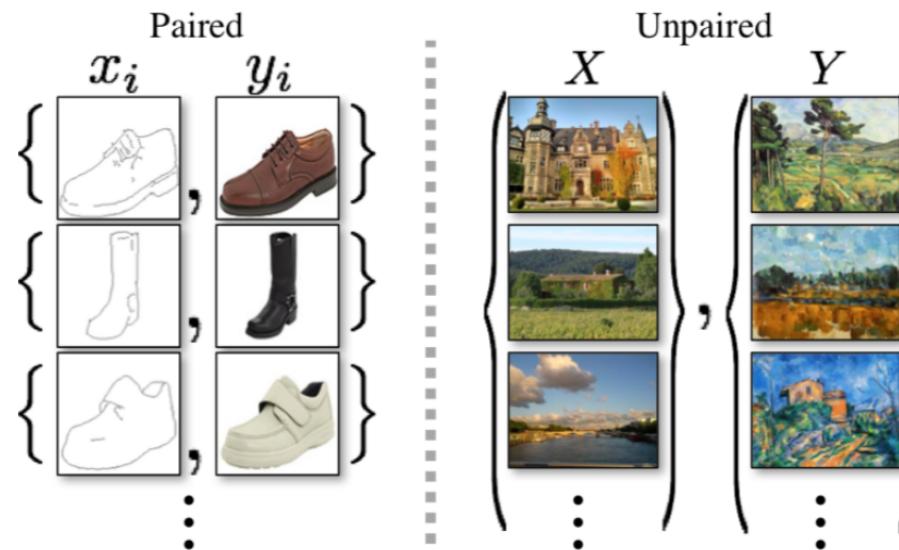


University at Buffalo The State University of New York

1846

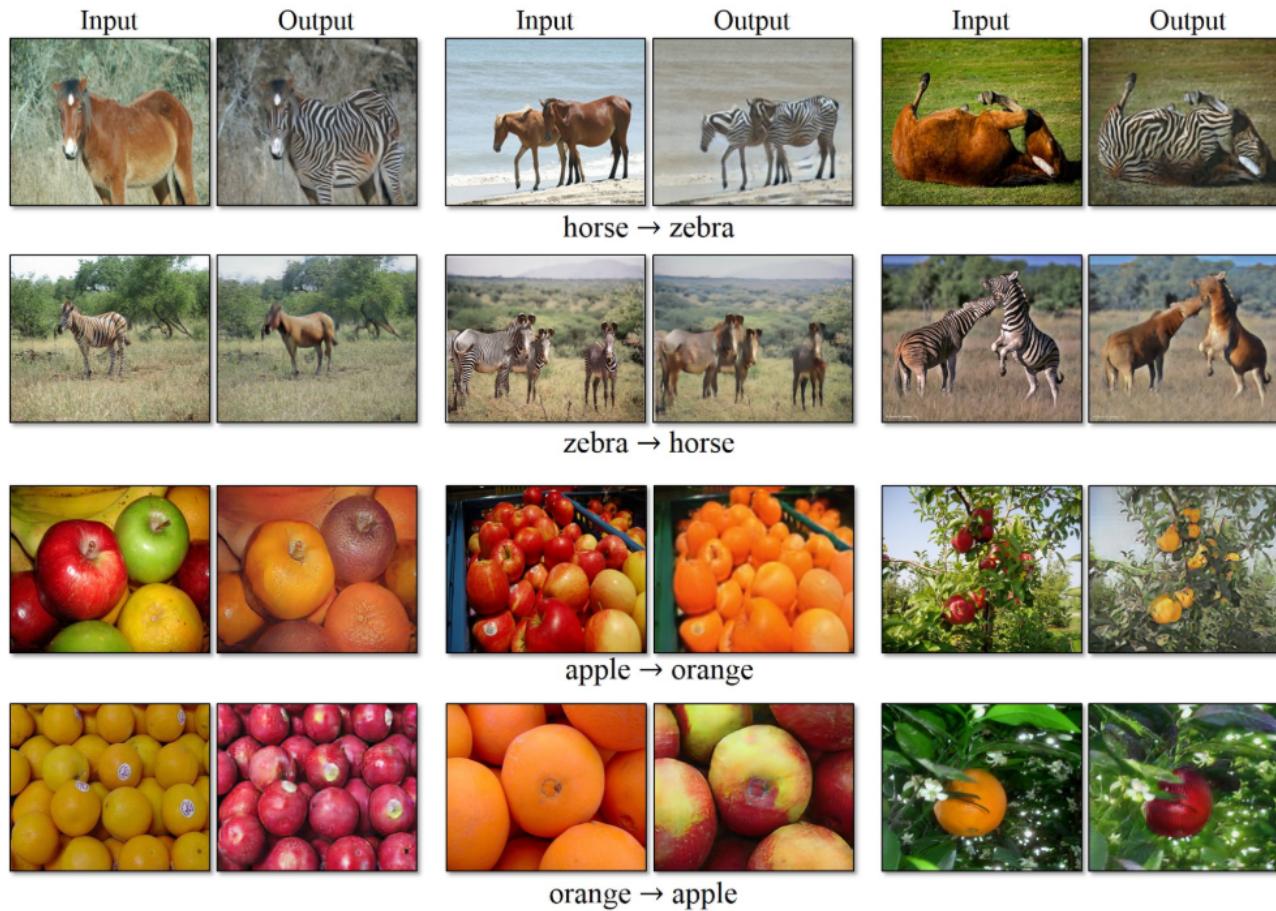
What is CycleGAN

- Instead of generating images from random noise, CycleGAN outputs different versions of input image.
- It is the first models to allow for *unpaired* image-to-image training.



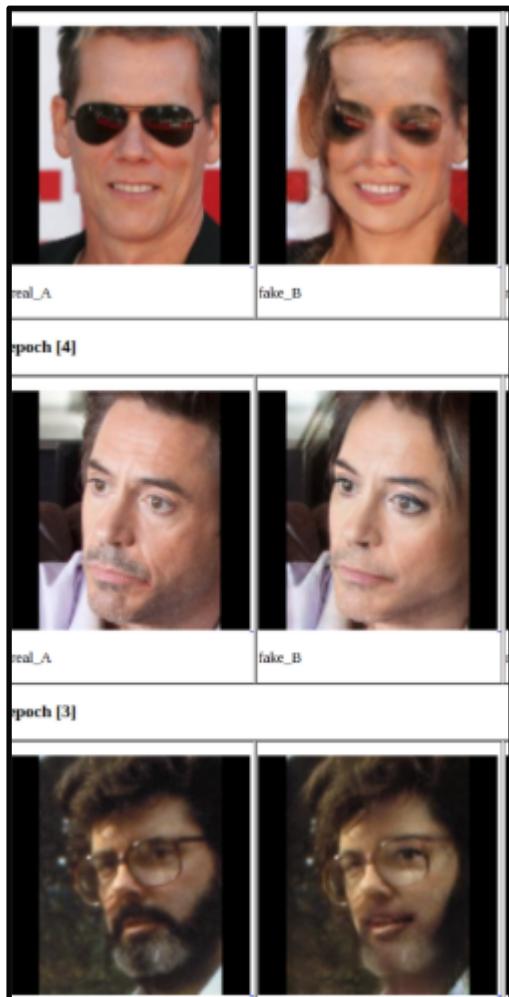
Zhu et al. 2017

CycleGAN Application



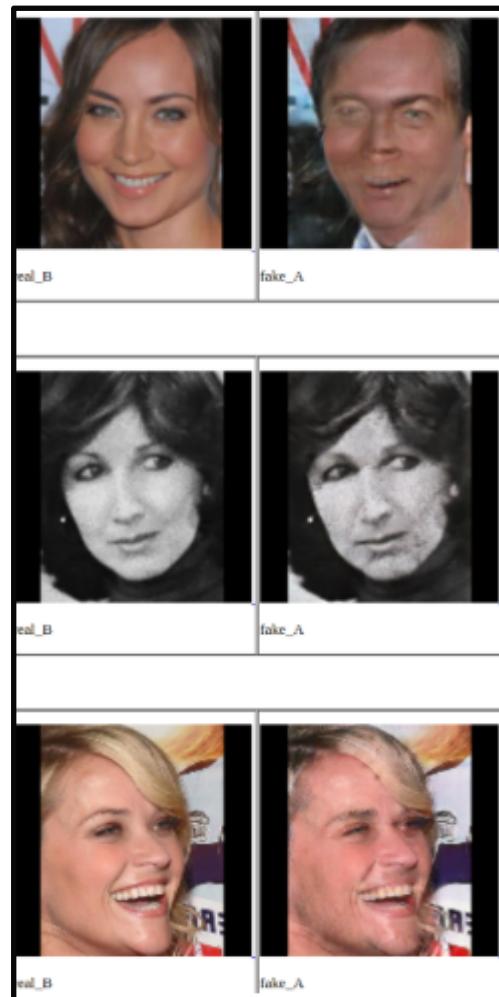
CycleGAN Application

Male → Female

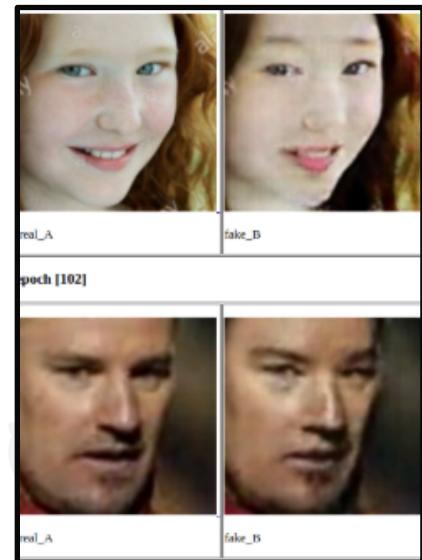


<http://shikib.com/CycleGan.html>

Female → Male



White → Asian



Asian → White



Introduction

- CycleGAN
- Applications

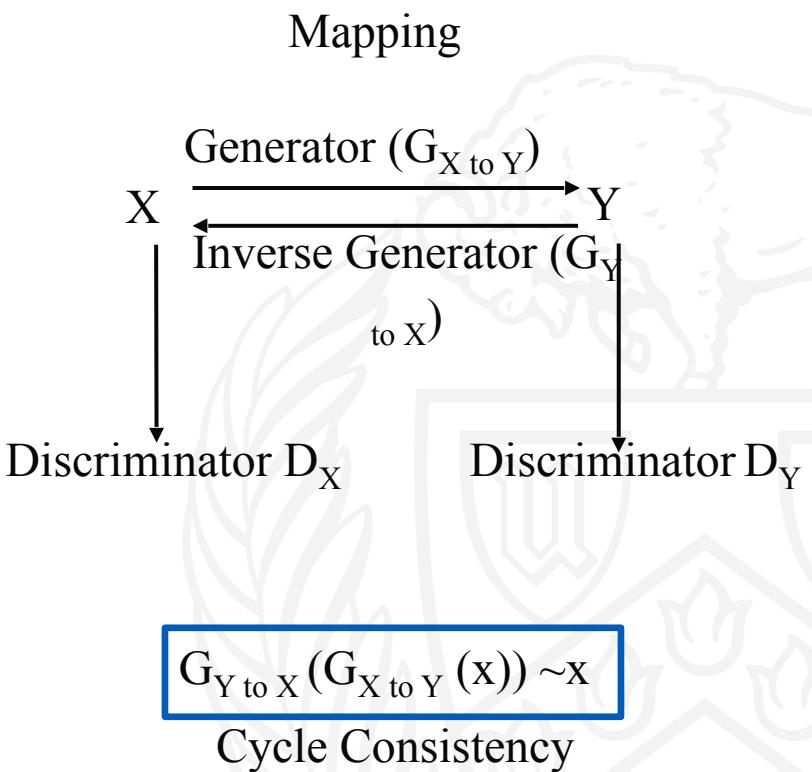
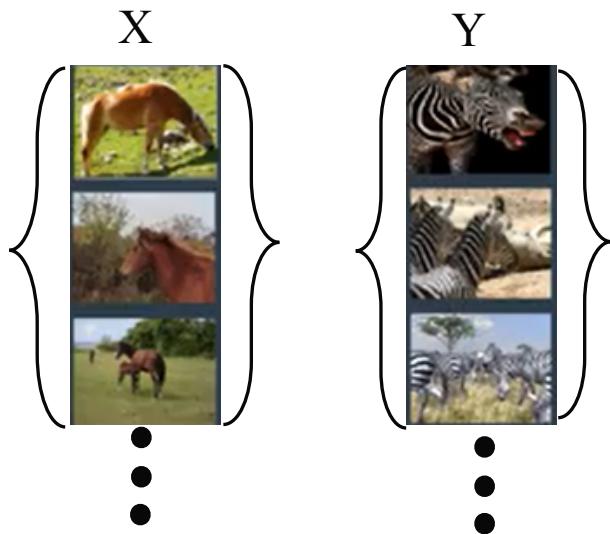
Methods

- Unpaired Data Mapping
- Model Architecture
- Cycle Consistency Loss
- Generator and Discriminator
- Loss functions
- Optimizer
- Model Training

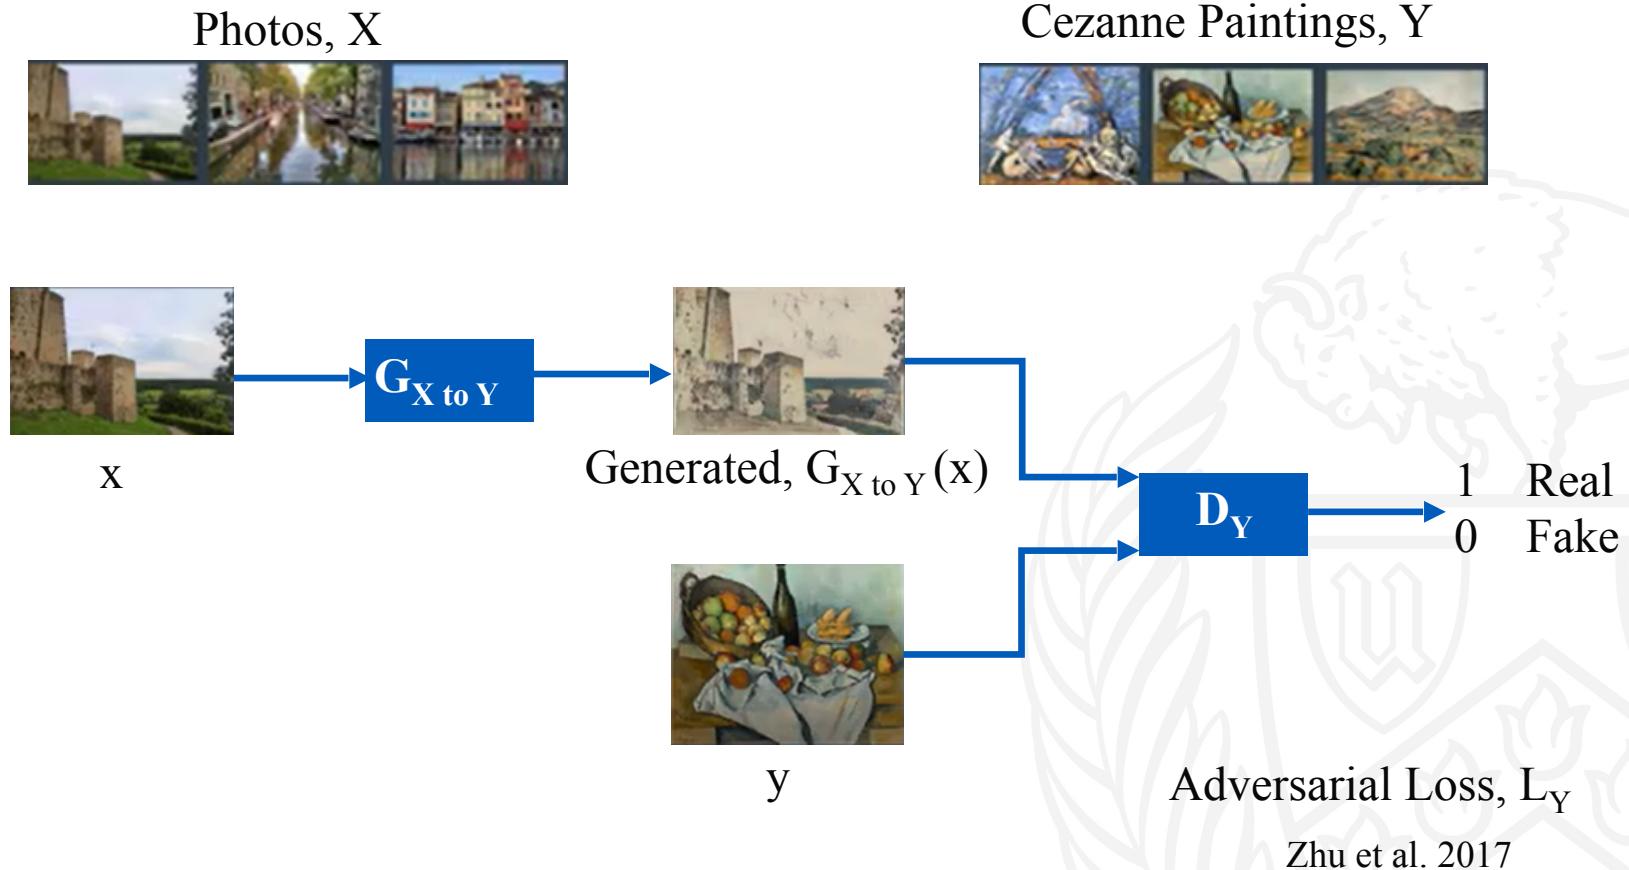
Discussion

- Broader Application
- Limitation
- Code Demo

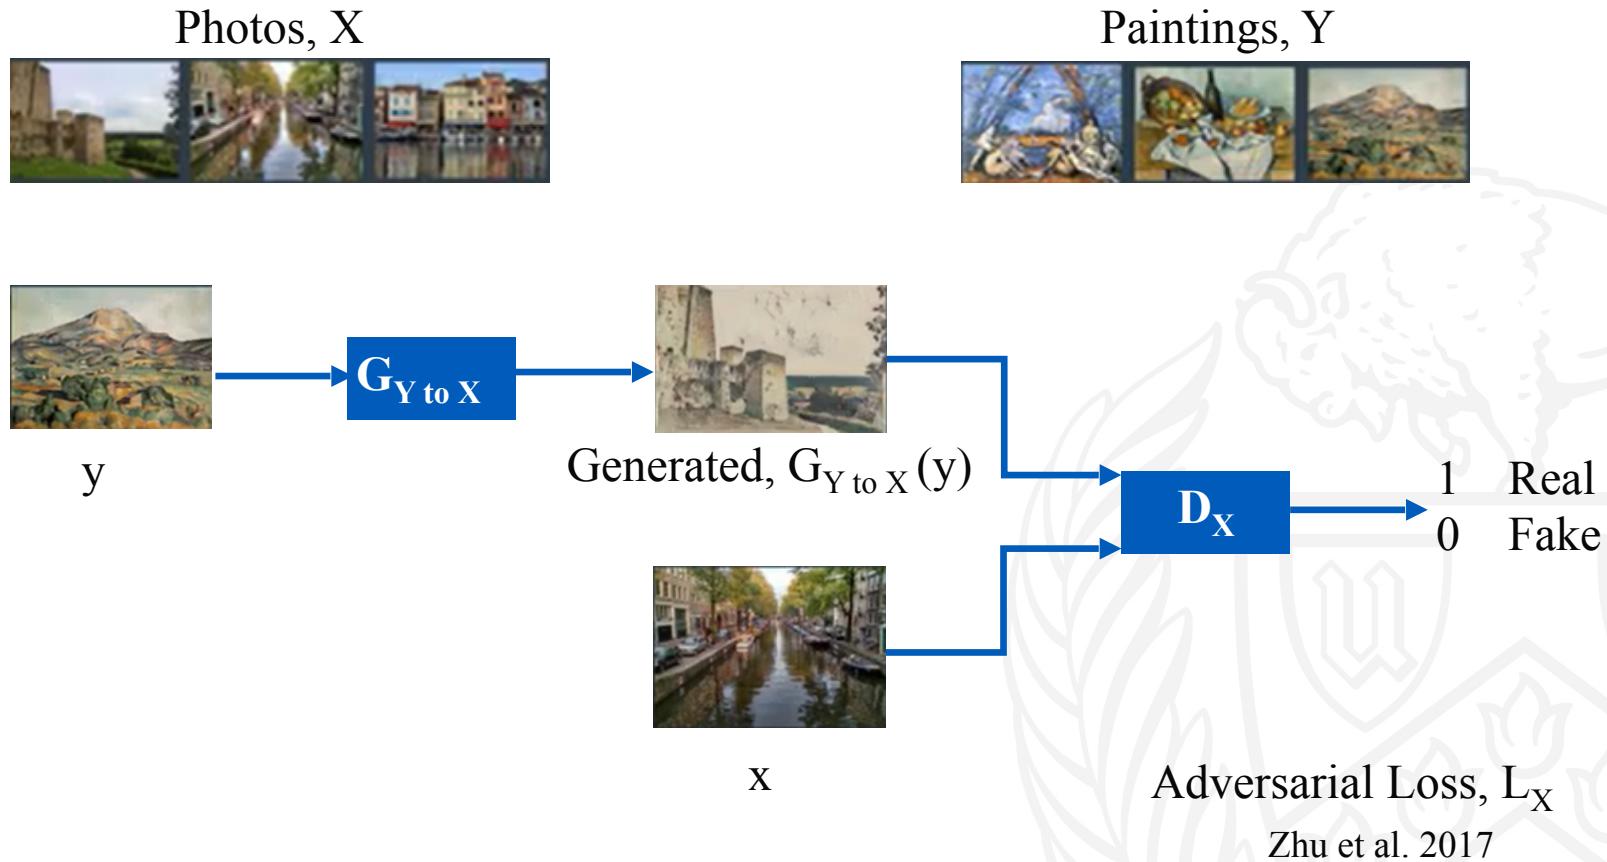
CycleGAN Unpaired Data Mapping



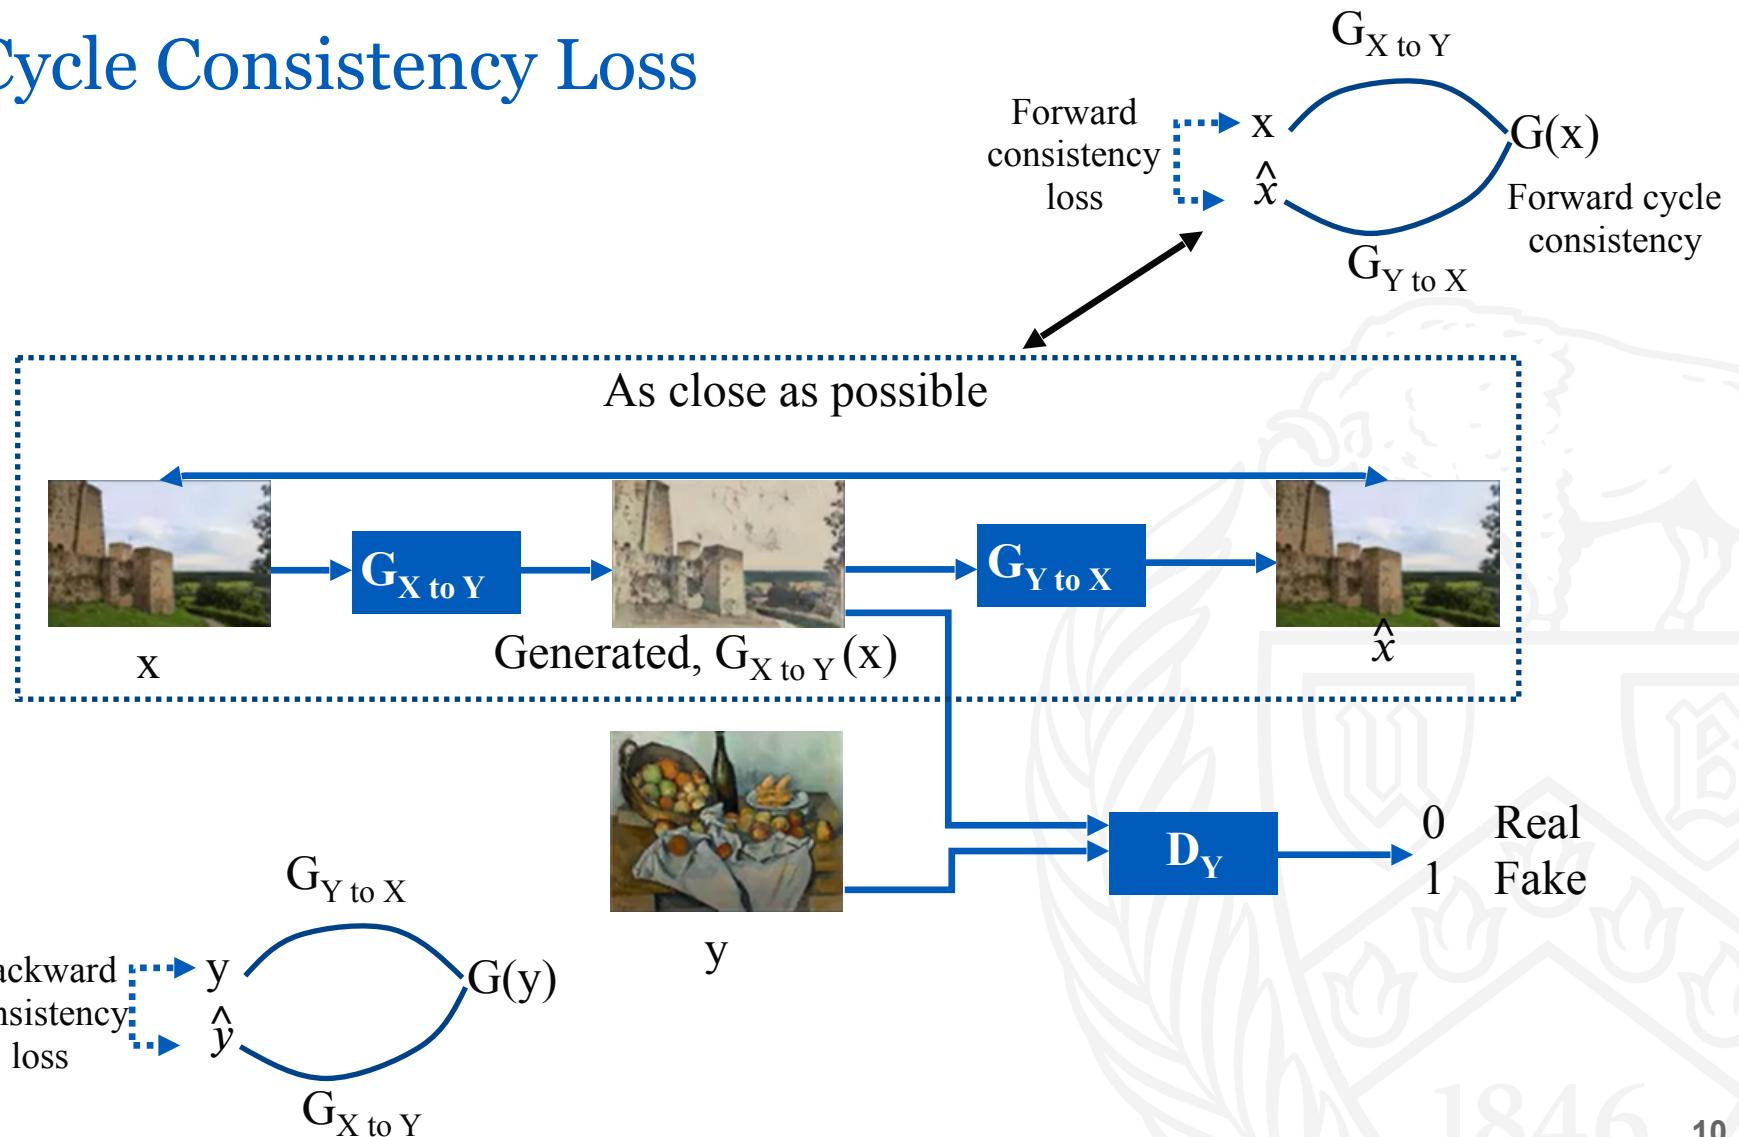
Model Architecture



Model Architecture



Cycle Consistency Loss



LOSS

- The full loss is adversarial loss + forward and backward cycle losses
- Adversarial loss is for style control
- Cycle consistency loss is for content control

$$L = L_y + L_x + \lambda L_{\text{cycle}}$$

→ Controls the relative importance of L_{cycle} .
Zhu et al. (2017) suggests 10.

Why Does CycleGAN work

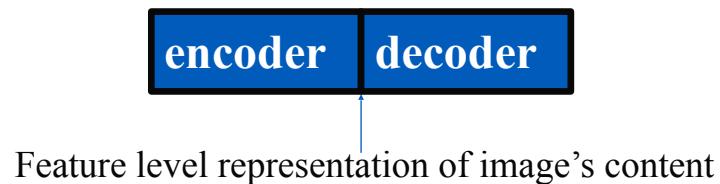


The assumption: the style and content of an image can be separated.

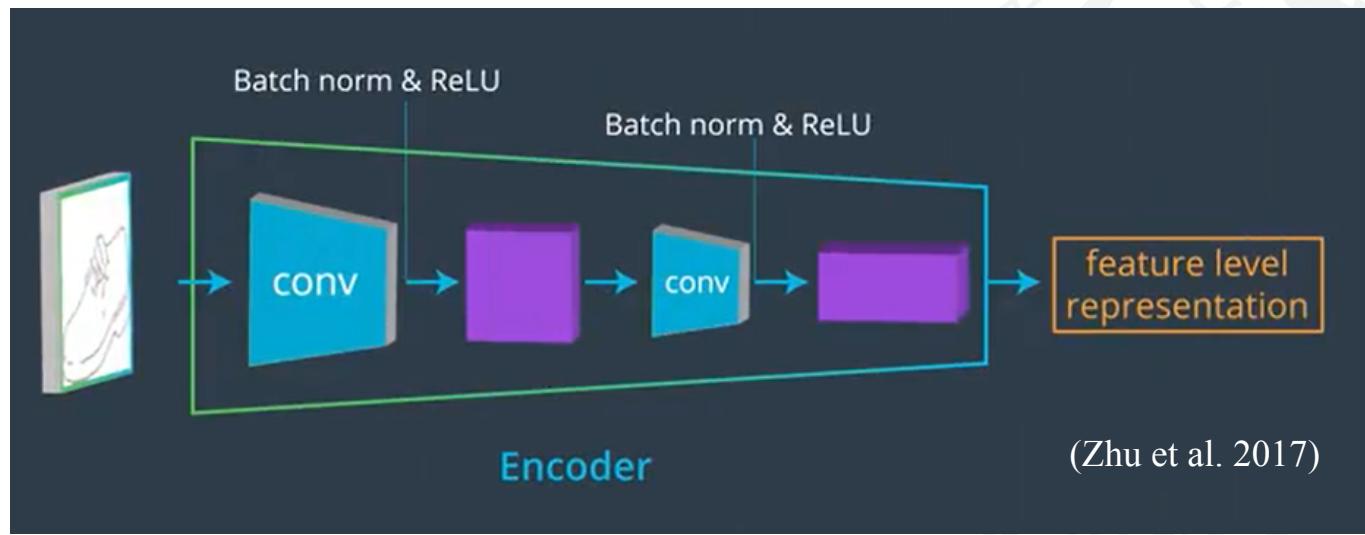


This is what the convolutional layers in the generators should learn

Generator ($G_{X \rightarrow Y}$ and $G_{Y \rightarrow X}$)

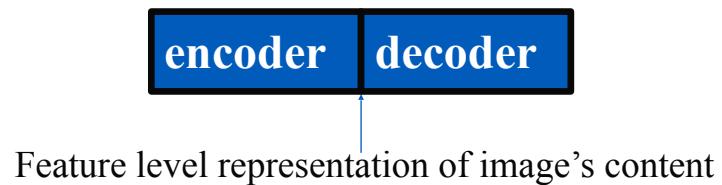


Encoder: convolutional neural network that extract features from the image with content and discard other information such as texture and colors.

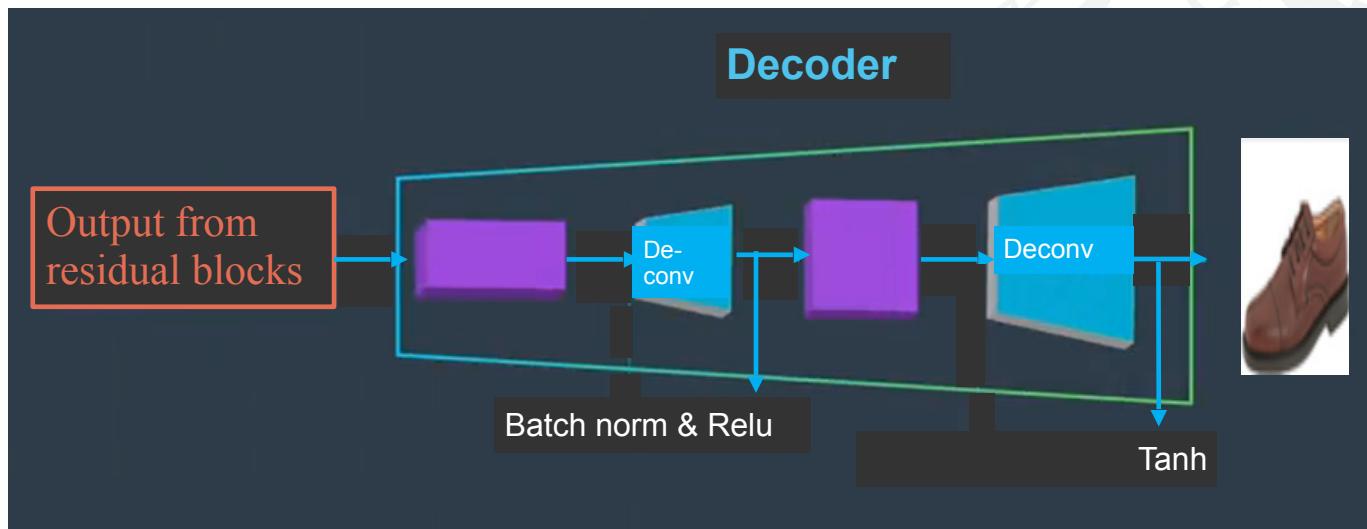


Layers increasing represent the content of the image

Generator ($G_{X \rightarrow Y}$ and $G_{Y \rightarrow X}$)



Decoder: transpose convolutional neural network that is responsible for turning the feature level representation into an transformed image



(Zhu et al. 2017)

Generator ($G_{X \rightarrow Y}$ and $G_{Y \rightarrow X}$)



Residual Blocks: connect the encoder and decoder

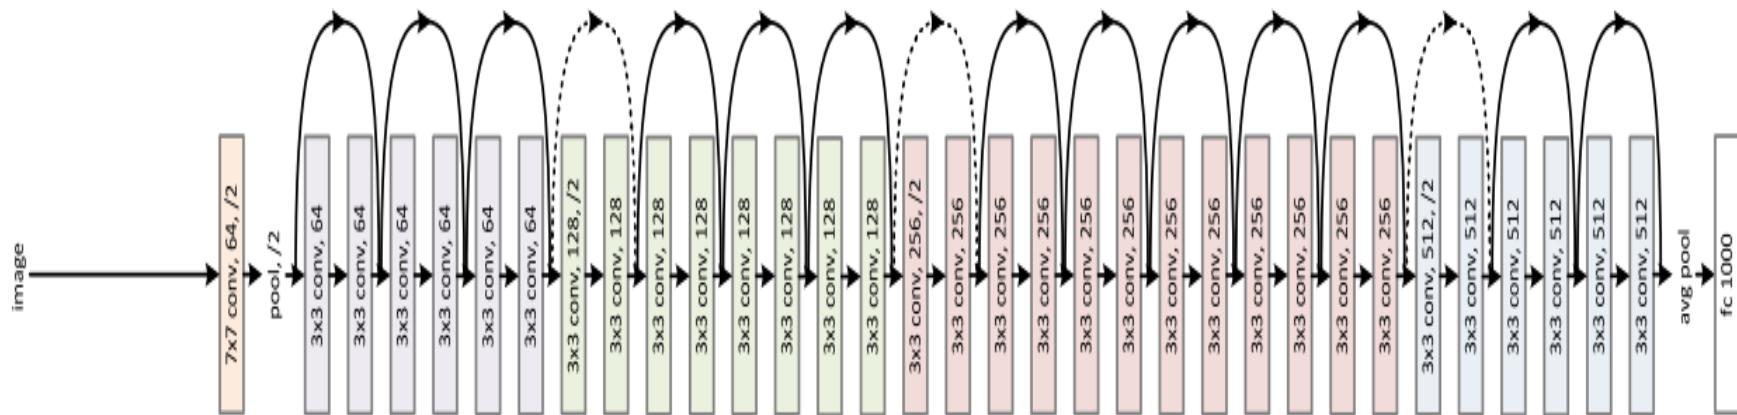
Why?

- Deeper networks are more likely to have vanishing or exploding gradients
- Have trouble reaching convergence.
- Batch normalization helps.
- Training accuracy stops improving at some point.
- In the worst cases, deep models would see their training accuracy worsen over time.

Generator ($G_{X \rightarrow Y}$ and $G_{Y \rightarrow X}$)



ResNet50 for image classification: connecting the output of one layer with the input of an earlier layer

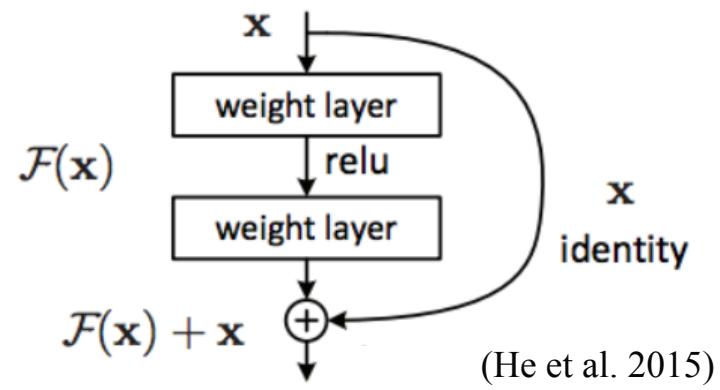


Generator ($G_{X \rightarrow Y}$ and $G_{Y \rightarrow X}$)

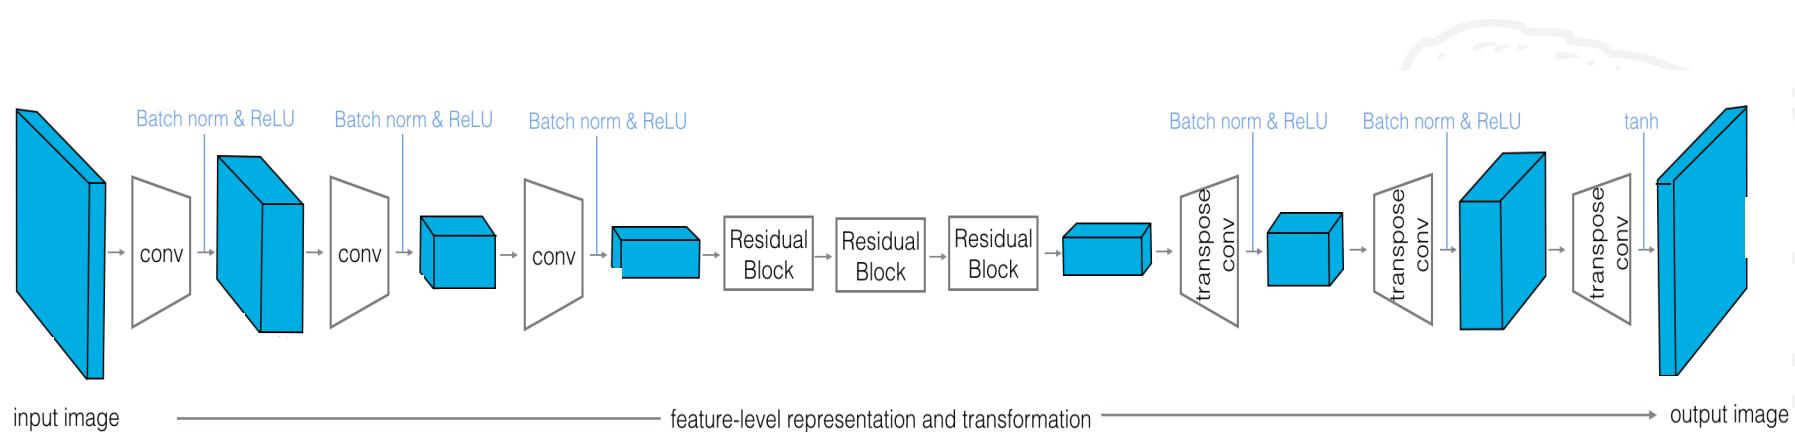


Residual blocks

- allow us to learn so-called *residual functions* $F(x)$
- typically, two convolutional layers + normalization layer and a ReLu in between.
- convolutional layers should have the same number of inputs as outputs.
- typically 6 or more blocks

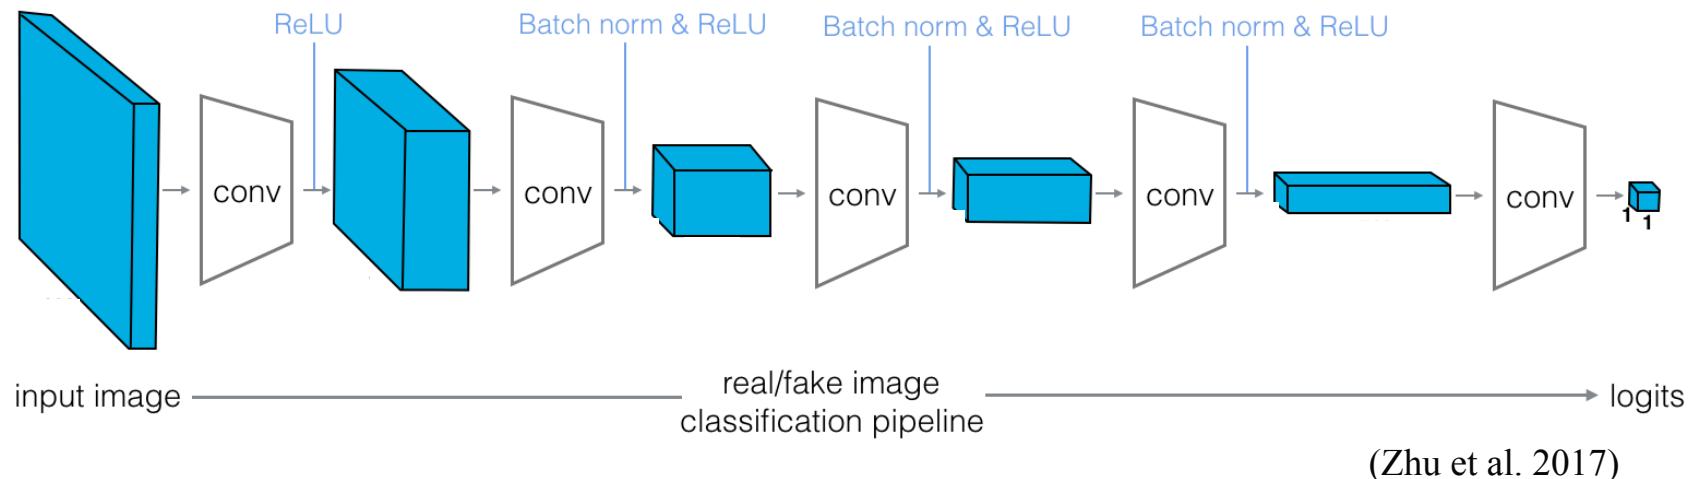


Generator ($G_{X \rightarrow Y}$ and $G_{Y \rightarrow X}$)



(Zhu et al. 2017)

Discriminator (D_X and D_Y)



More About Loss Functions

The full loss is adversarial loss + forward and backward cycle losses

Adversarial loss:

$$L_{DY} = E_{y \sim pdata(y)} [(D_Y(y) - 1)^2] + E_{x \sim pdata(x)} [D_Y(G(x))^2]$$

real loss fake loss

Discriminator

$$L_{DX} = E_{x \sim pdata(x)} [(D_X(x) - 1)^2] + E_{y \sim pdata(y)} [D_X(F(y))^2]$$

real loss fake loss

Generator

$$L_G = E_{x \sim p_{\text{data}}(x)} [D_Y (G(x) - 1)^2]$$

$$L_F = E_{y \sim p(\text{data}|y)} [D_X (F(y) - 1)^2]$$

Least-square loss is more
stable/gradient vanish

More About Loss Functions

The full loss is adversarial loss + forward and backward cycle losses

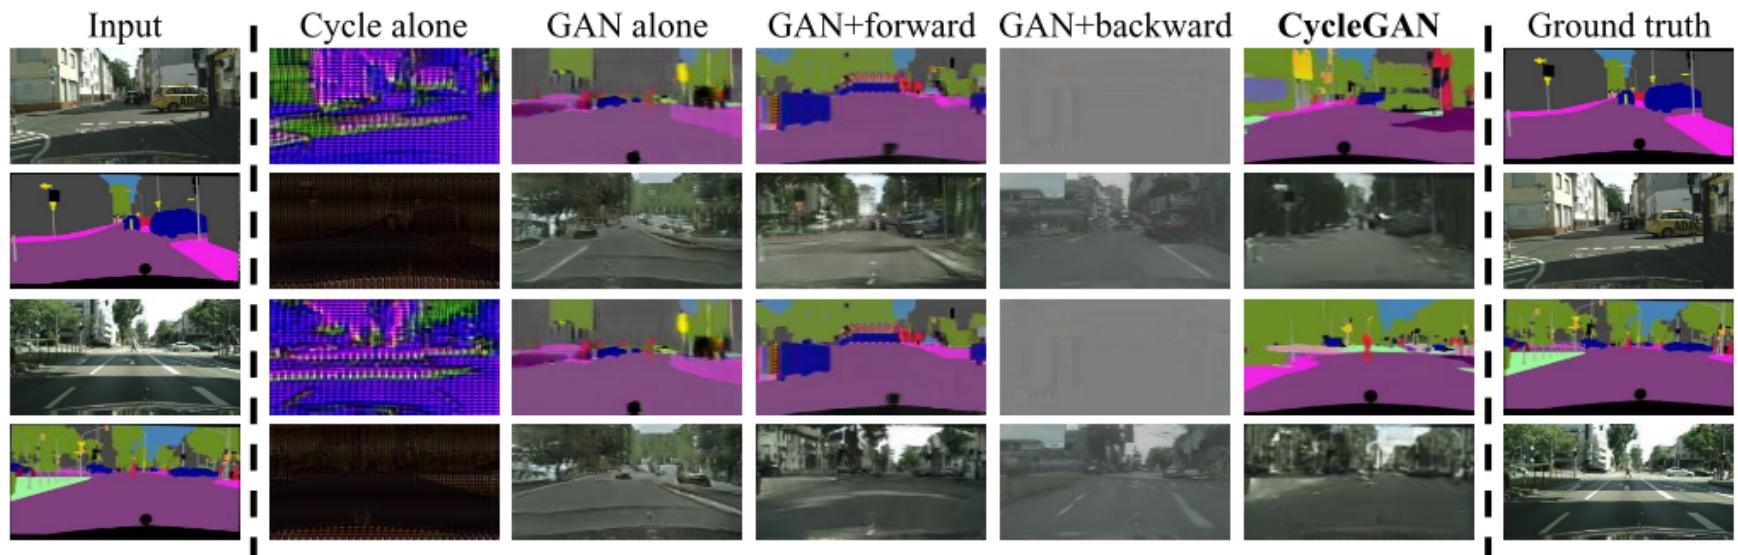
Cycle consistency loss:

$$L_{cycle} = E_{x \sim p_{data}(x)} [\|F(G(x)) - x\|_1] + E_{y \sim p_{data}(y)} [\|G(F(y)) - y\|_1]$$

Forward consistency loss Backward consistency loss L1 norm

(Zhu et al. 2017)

More About Loss Functions



(Zhu et al. 2017)

Optimizer

Adam with adapted learning rate:

- all networks are trained from scratch with a learning rate of 0.0002 for the first 100 epochs
- linearly decay the rate to zero over the next 100 epochs

(Zhu et al. 2017)

1846

23

General Training Steps

Training the Discriminators

1. Compute the discriminator D_X loss on real images
2. Generate fake images that look like domain X based on real images in domain Y
3. Compute the fake loss for D_X
4. Compute the total loss and perform backpropagation and D_X optimization
5. Repeat steps 1-4 only with D_Y and your domains switched!

Training the Generators

1. Generate fake images that look like domain X based on real images in domain Y
2. Compute the generator loss based on how D_X responds to fake X
3. Generate reconstructed \hat{Y} images based on the fake X images generated in step 1
4. Compute the cycle consistency loss by comparing the reconstructions with real Y images
5. Repeat steps 1-4 only swapping domains
6. Add up all the generator and reconstruction losses and perform backpropagation + optimization

(Zhu et al. 2017)

Introduction

- CycleGAN
- Applications

Methods

- Unpaired Data Mapping
- Model Architecture
- Cycle Consistency Loss
- Generator and Discriminator
- Loss functions
- Optimizer
- Model Training

Discussion

- Broader Application
- Limitation
- Code Demo

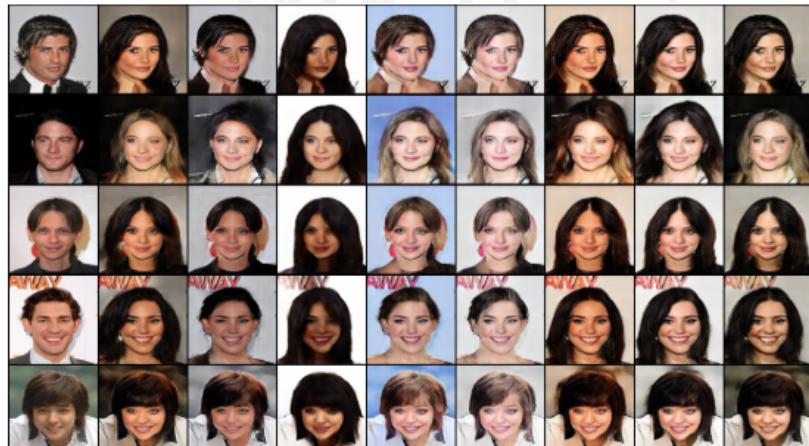


University at Buffalo The State University of New York

1846 25

Boarder Applications

- [Video-to-Video](#) (Zhu et al. 2017. Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks)
- [Voice-to-Voice](#) (Kaneko and Kameoka 2017. Parallel-Data-Free Voice Conversion Using Cycle-Consistent Adversarial Networks)
- [Faces-off](#) (Jin et al. 2018. CycleGAN Face-off)
- Many-to-Many Mappings (Almahairi et al. 2018. Augmented CycleGAN: Learning Many-to-Many Mappings from Unpaired Data)



Shortcomings

- It will only show one version of a transformed output even if there are multiple possible outputs.
- A simple CycleGAN produces low-resolution images, though there is some research around high-resolution GANs (Wang et al. 2018. High-Resolution Image Synthesis and Semantic Manipulation with Conditional GANs)
- Good at color and texture changes, while bad at geometric changes

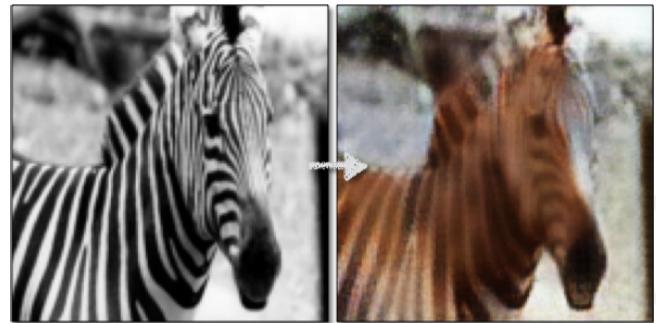


(Zhu et al. 2017)

Some Improvements

- Cycle consistency on discriminator CNN feature level

Modified cycle consistency loss as a linear combination of discriminator feature extractor (last layer) and pixel level consistency losses



$$\tilde{\mathcal{L}}_{\text{cyc}}(G, F, D_X, X, \gamma) = \mathbb{E}_{x \sim p_{\text{data}}(x)} [\gamma \|f_{D_X}(F(G(x))) - f_{D_X}(x)\|_1 + (1-\gamma) \|F(G(x)) - x\|_1]$$

Wang and Lin 2018

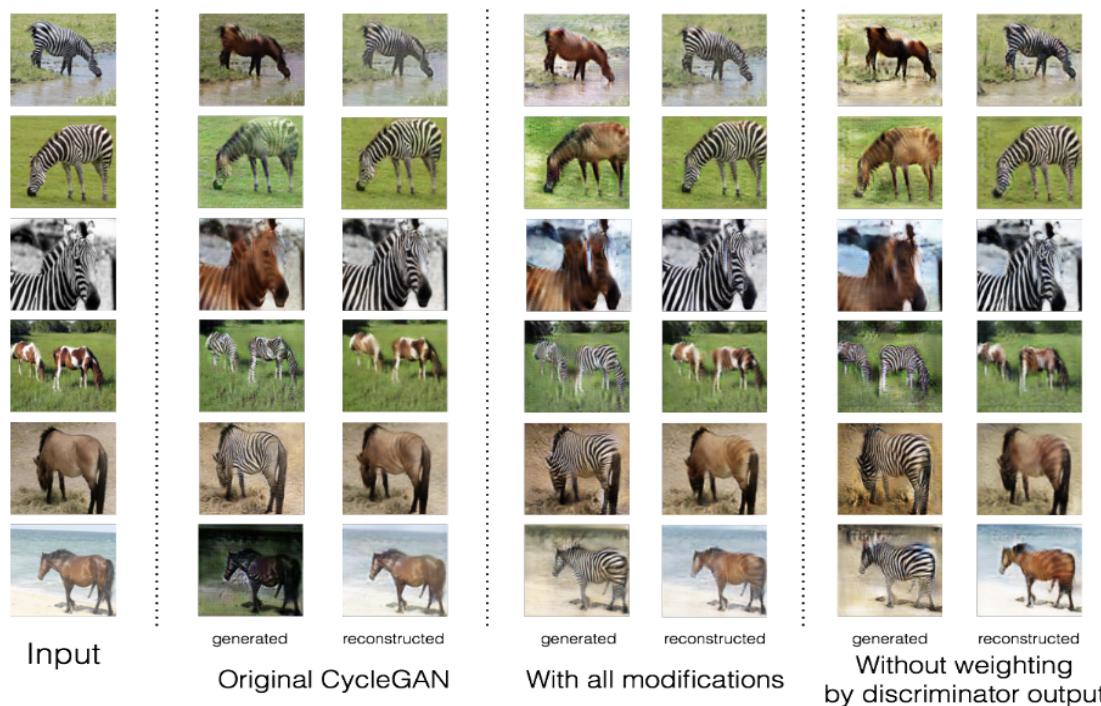
- Cycle consistency weight decay

Gradually decay the weight of cycle consistency loss λ as training progress but keep it bigger than 0.

Some Improvements

- Weight cycle consistency by quality of generated image

$$\tilde{\mathcal{L}}_{\text{cyc}}(G, F, D_X, X, \gamma) = \mathbb{E}_{x \sim p_{\text{data}}(x)} \left[D_X(x) \left(\gamma \|f_{D_X}(F(G(x))) - f_{D_X}(x)\|_1 + (1 - \gamma) \|F(G(x)) - x\|_1 \right) \right]$$



Code Demo



THANK YOU FOR YOUR ATTENTION!

Q & A!

